

CS457 – P2P

Fall 2014

Topics

- Types of peer-to-peer networks
 - Directory-based (e.g., original Napster design)
 - Unstructured (e.g., Gnutella, Kazaa, BitTorrent)
 - Structured (e.g., distributed hash tables) (later)
- Challenges in peer-to-peer
 - Legal issues, free riding, fast response to queries, peers coming and going over time, reliability, security, ...

Peer-to-Peer Networks: Napster



**Shawn Fanning,
Northeastern freshman**

- Napster history: the rise
 - January 1999: Napster version 1.0
 - May 1999: company founded
 - September 1999: first lawsuits
 - 2000: 80 million users
- Napster history: the fall
 - Mid 2001: out of business due to lawsuits
 - Mid 2001: dozens of P2P alternatives that were harder to touch, though these have gradually been constrained
 - 2003: growth of pay services like iTunes
- Napster history: the resurrection
 - 2003: Napster reconstituted as a pay service
 - 2006: still lots of file sharing going on

Napster Technology: Directory Service



- User installs the software
 - Download the client program
 - Register name, password, local directory, etc.
- Client contacts Napster (via TCP)
 - Provides a list of music files it will share
 - ... and Napster's central server updates the directory
- Client searches on a title or performer
 - Napster identifies online clients with the file
 - ... and provides IP addresses
- Client requests the file from the chosen supplier
 - Supplier transmits the file to the client
 - Both client and supplier report status to Napster

Napster Technology: Properties

- Server's directory continually updated
 - Always know what music is currently available
 - Point of vulnerability for legal action
- Peer-to-peer file transfer
 - No load on the server
 - Plausible deniability for legal action (but not enough)
- Proprietary protocol
 - Login, search, upload, download, and status operations
 - No security: clear-text passwords and other vulnerabilities
- Bandwidth issues
 - Suppliers ranked by apparent bandwidth & response time

Napster: Limitations of Central Directory

- Single point of failure
- Performance bottleneck
- Copyright infringement

File transfer is decentralized, but locating content is highly centralized

- So, later P2P systems were more distributed

Peer-to-Peer Networks: Gnutella

- Gnutella history
 - 2000: J. Frankel & T. Pepper released Gnutella
 - Soon after: many other clients (e.g., Morpheus, Limewire, Bearshare)
 - 2001: protocol enhancements, e.g., “ultrapeers”
- Query flooding
 - Join: contact a few nodes to become neighbors
 - Publish: no need!
 - Search: ask neighbors, who ask their neighbors
 - Fetch: get file directly from another node



Gnutella: Query Flooding

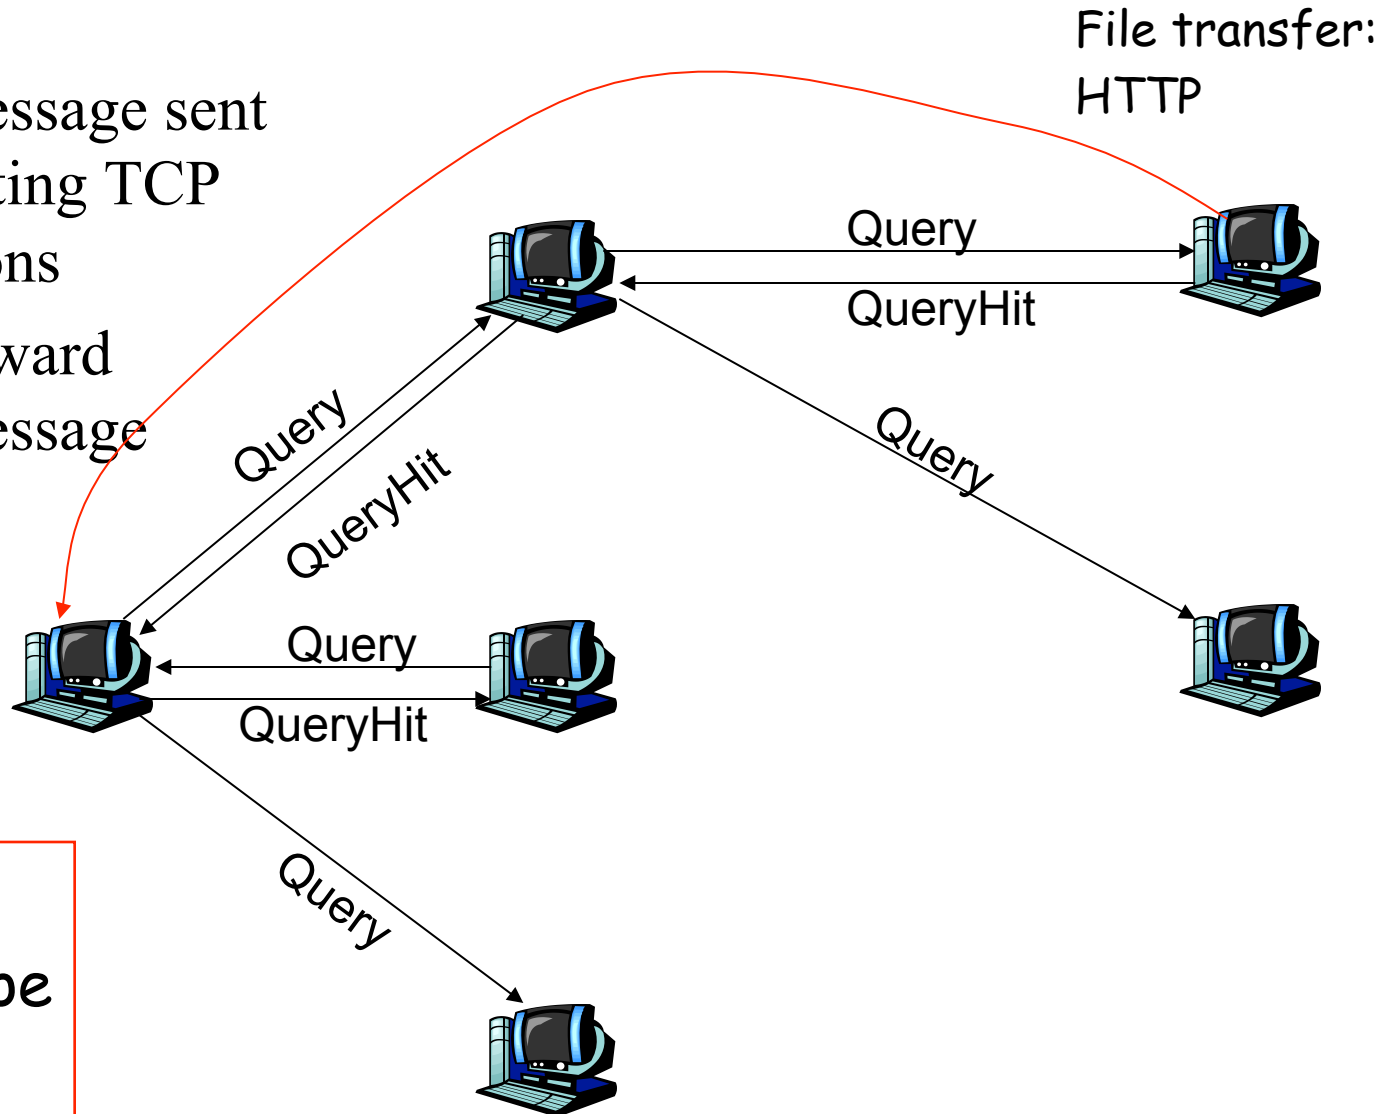
- Fully distributed
 - No central server
- Public domain protocol
- Many Gnutella clients implement the protocol

Overlay network: graph

- Edge between peer X and Y if there's a TCP connection
- All active peers and edges form the overlay network
- Given peer will typically be connected with < 10 overlay neighbors

Gnutella: Protocol

- Query message sent over existing TCP connections
- Peers forward Query message
- QueryHit sent over reverse path



Scalability:
limited scope
flooding

Gnutella: Peer Joining

- Joining peer X must find some other peer in Gnutella network: use list of candidate peers
- X sequentially attempts to make TCP connection with peers on list until connection setup with Y
- X sends Ping message to Y; Y forwards Ping message.
- Peers receiving Ping message may respond with Pong message
- X receives many Pong messages. It can then setup additional TCP connections

Gnutella: Pros and Cons

- Advantages
 - Fully decentralized
 - Search cost distributed
 - Processing per node permits powerful search semantics
- Disadvantages
 - Search scope may be quite large
 - Search time may be quite long
 - High overhead and nodes come and go often


Peer-to-Peer Networks: KaAzA

- KaZaA history
 - 2001: created by Dutch company (Kazaa BV)
 - Single network called FastTrack used by other clients as well
 - Eventually the protocol changed so other clients could no longer talk to it
- Smart query flooding
 - Join: on start, the client contacts a super-node (and may later become one)
 - Publish: client sends list of files to its super-node
 - Search: send query to super-node, and the super-nodes flood queries among themselves
 - Fetch: get file directly from peer(s); can fetch from multiple peers at once



KaZaA: Exploiting Heterogeneity

- Each peer is either a group leader or assigned to a group leader
 - TCP connection between peer and its group leader
 - TCP connections between some pairs of group leaders
- Group leader tracks the content in all its children



KaZaA: Motivation for Super-Nodes

- Query consolidation
 - Many connected nodes may have only a few files
 - Propagating query to a sub-node may take more time than for the super-node to answer itself
- Stability
 - Super-node selection favors nodes with high up-time
 - How long you've been up is a good predictor of how long you'll be around in the future

Peer-to-Peer Networks:

BitTorrent

- BitTorrent history and motivation
 - 2002: B. Cohen debuted BitTorrent
 - Key motivation: popular content
 - Popularity exhibits temporal locality (Flash Crowds)
 - E.g., Slashdot effect, CNN Web site on 9/11, release of a new movie or game
 - Focused on efficient *fetching*, not searching
 - Distribute same file to many peers
 - Single publisher, many downloaders
 - Preventing free-loading



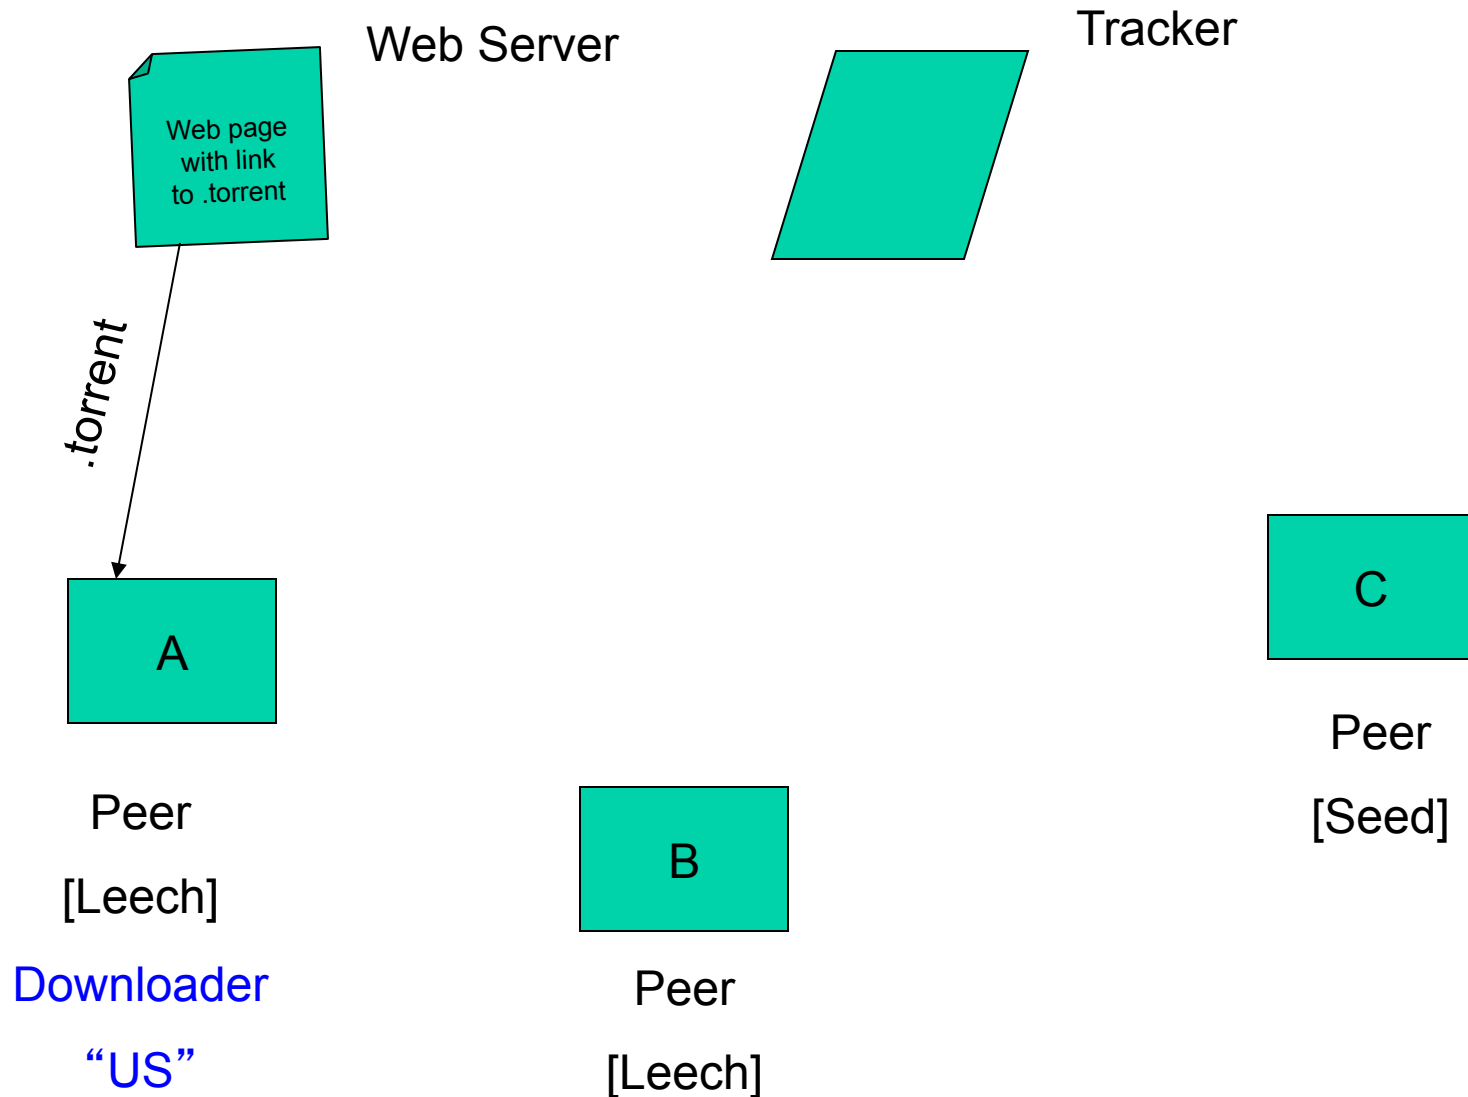
BitTorrent: Simultaneous Downloading

- Divide large file into many pieces
 - Replicate different pieces on different peers
 - A peer with a complete piece can trade with other peers
 - Peer can (hopefully) assemble the entire file
- Allows simultaneous downloading
 - Retrieving different parts of the file from different peers at the same time

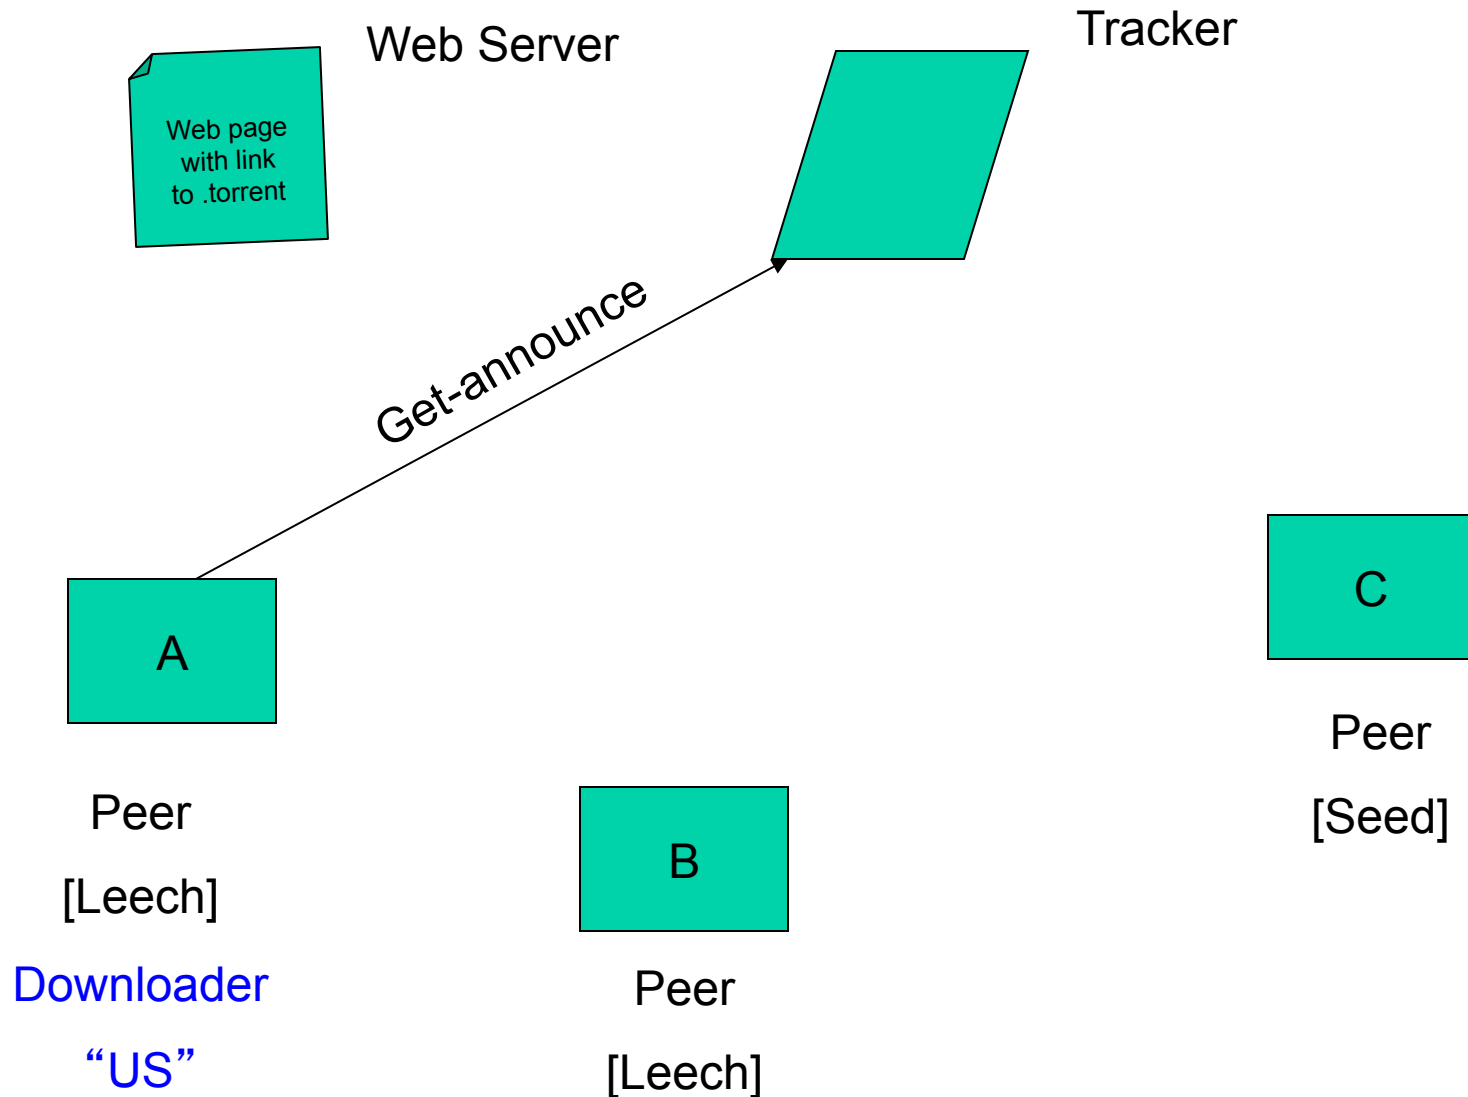
BitTorrent Components

- Seed
 - Peer with entire file
 - Fragmented in pieces
- Leacher
 - Peer with an incomplete copy of the file
- Torrent file
 - Passive component
 - Stores summaries of the pieces to allow peers to verify their integrity
- Tracker
 - Allows peers to find each other
 - Returns a list of random peers

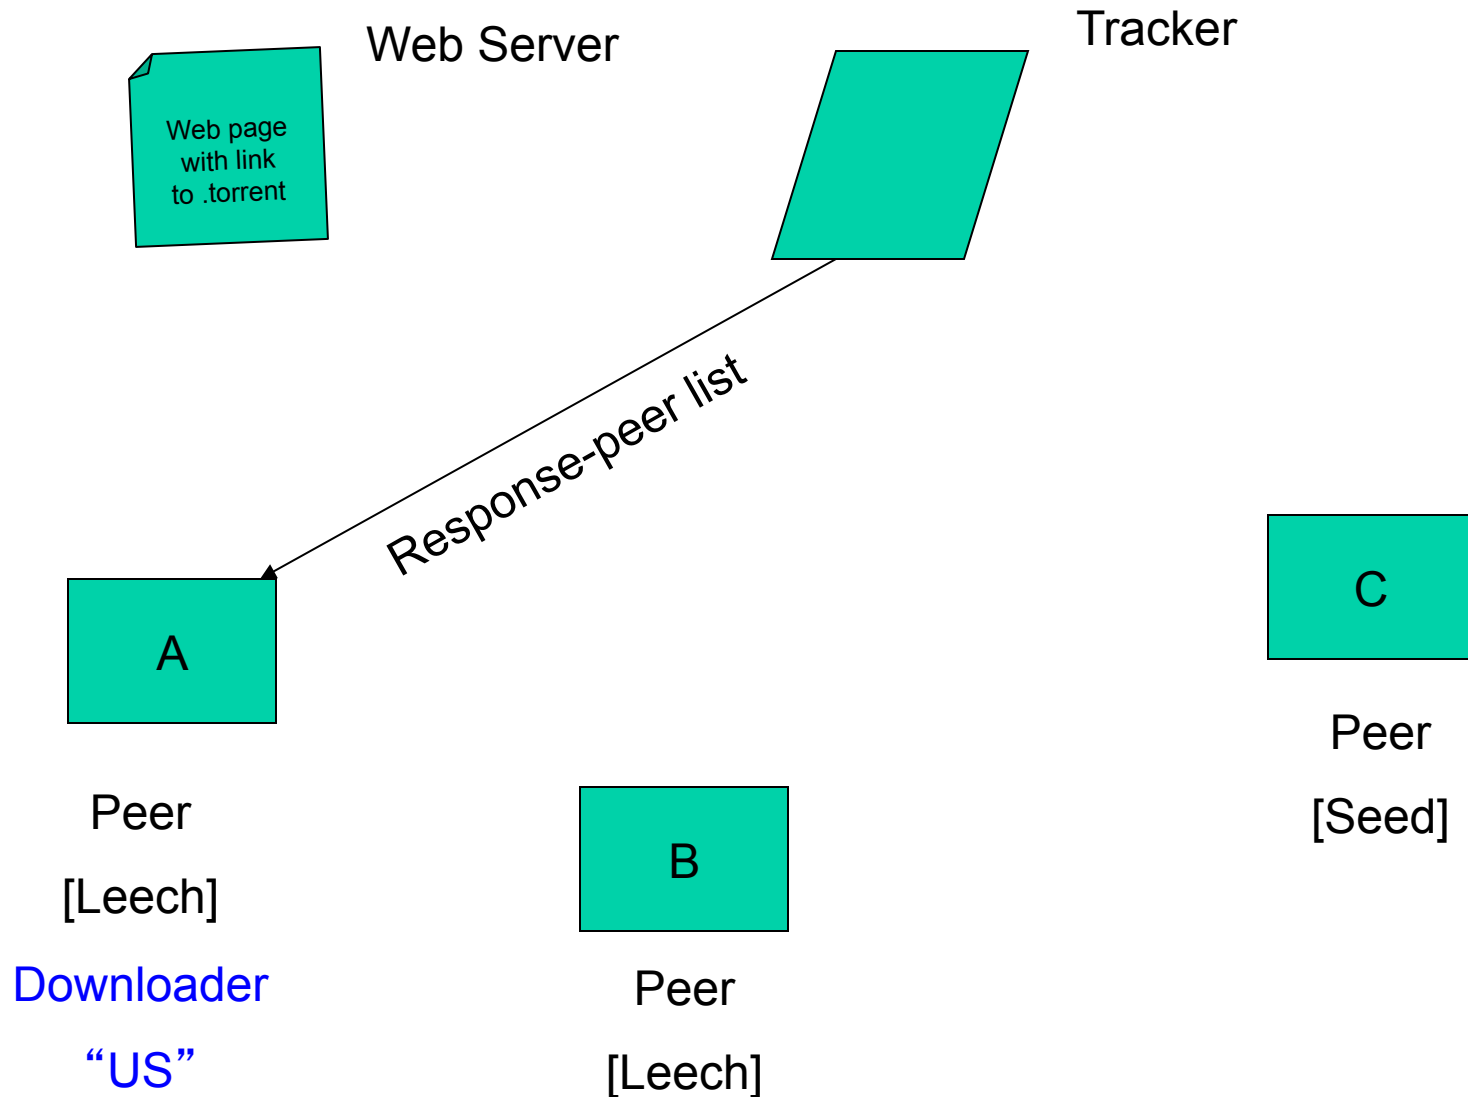
BitTorrent: Overall Architecture



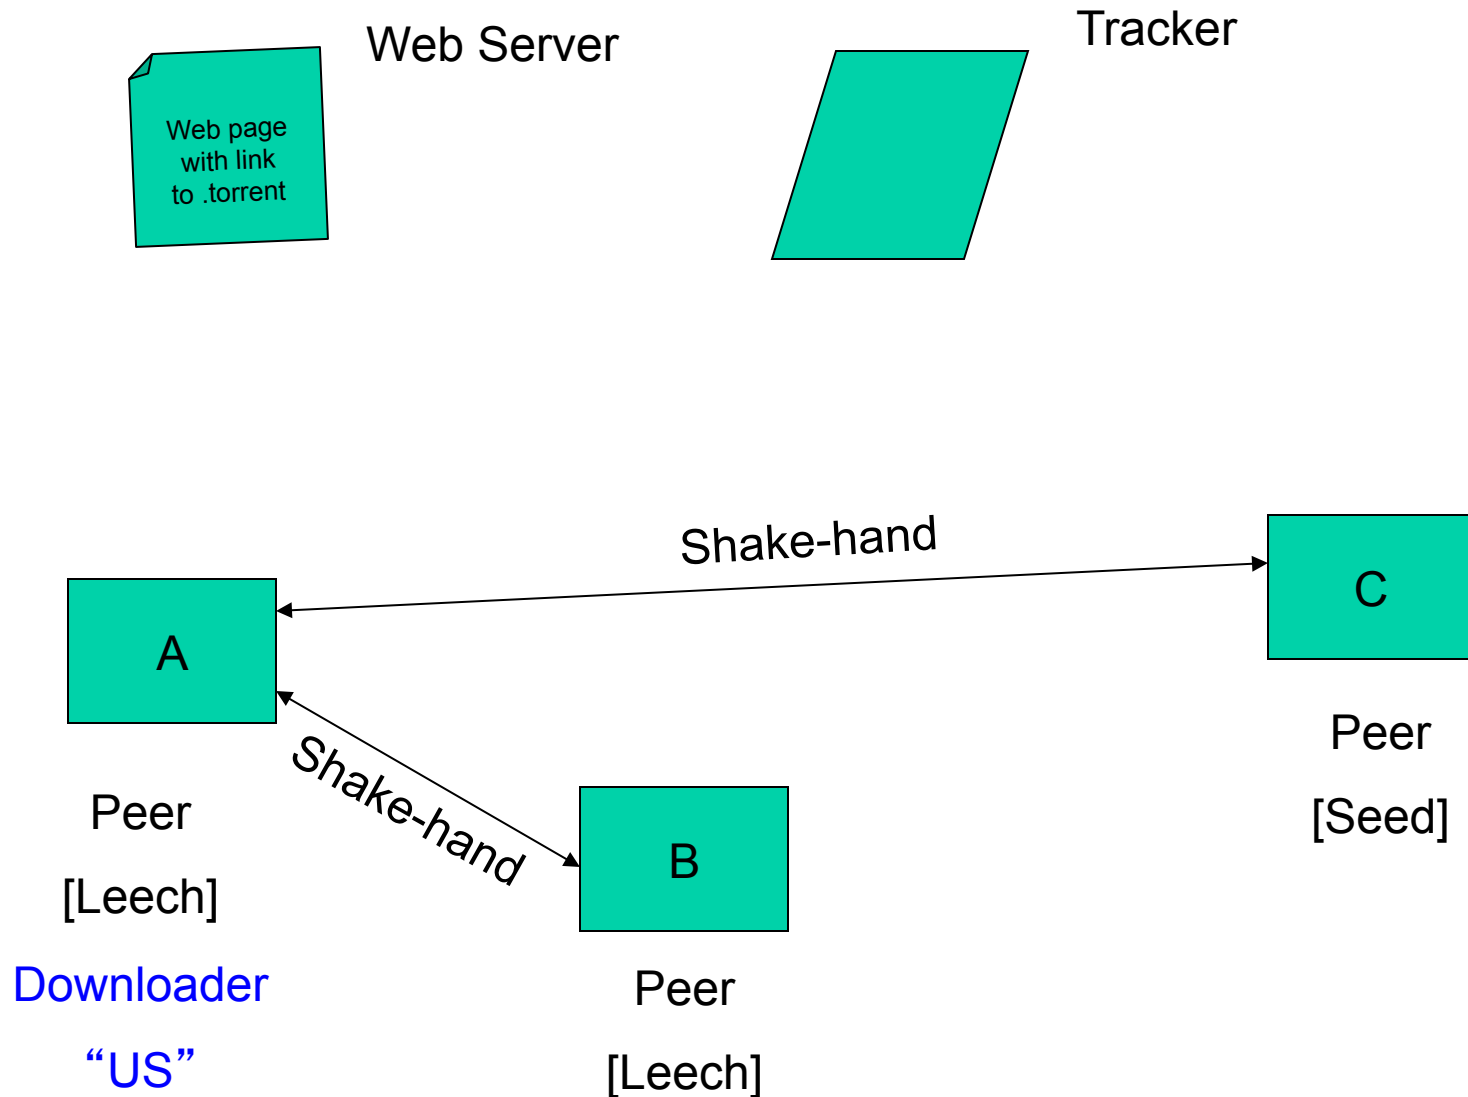
BitTorrent: Overall Architecture



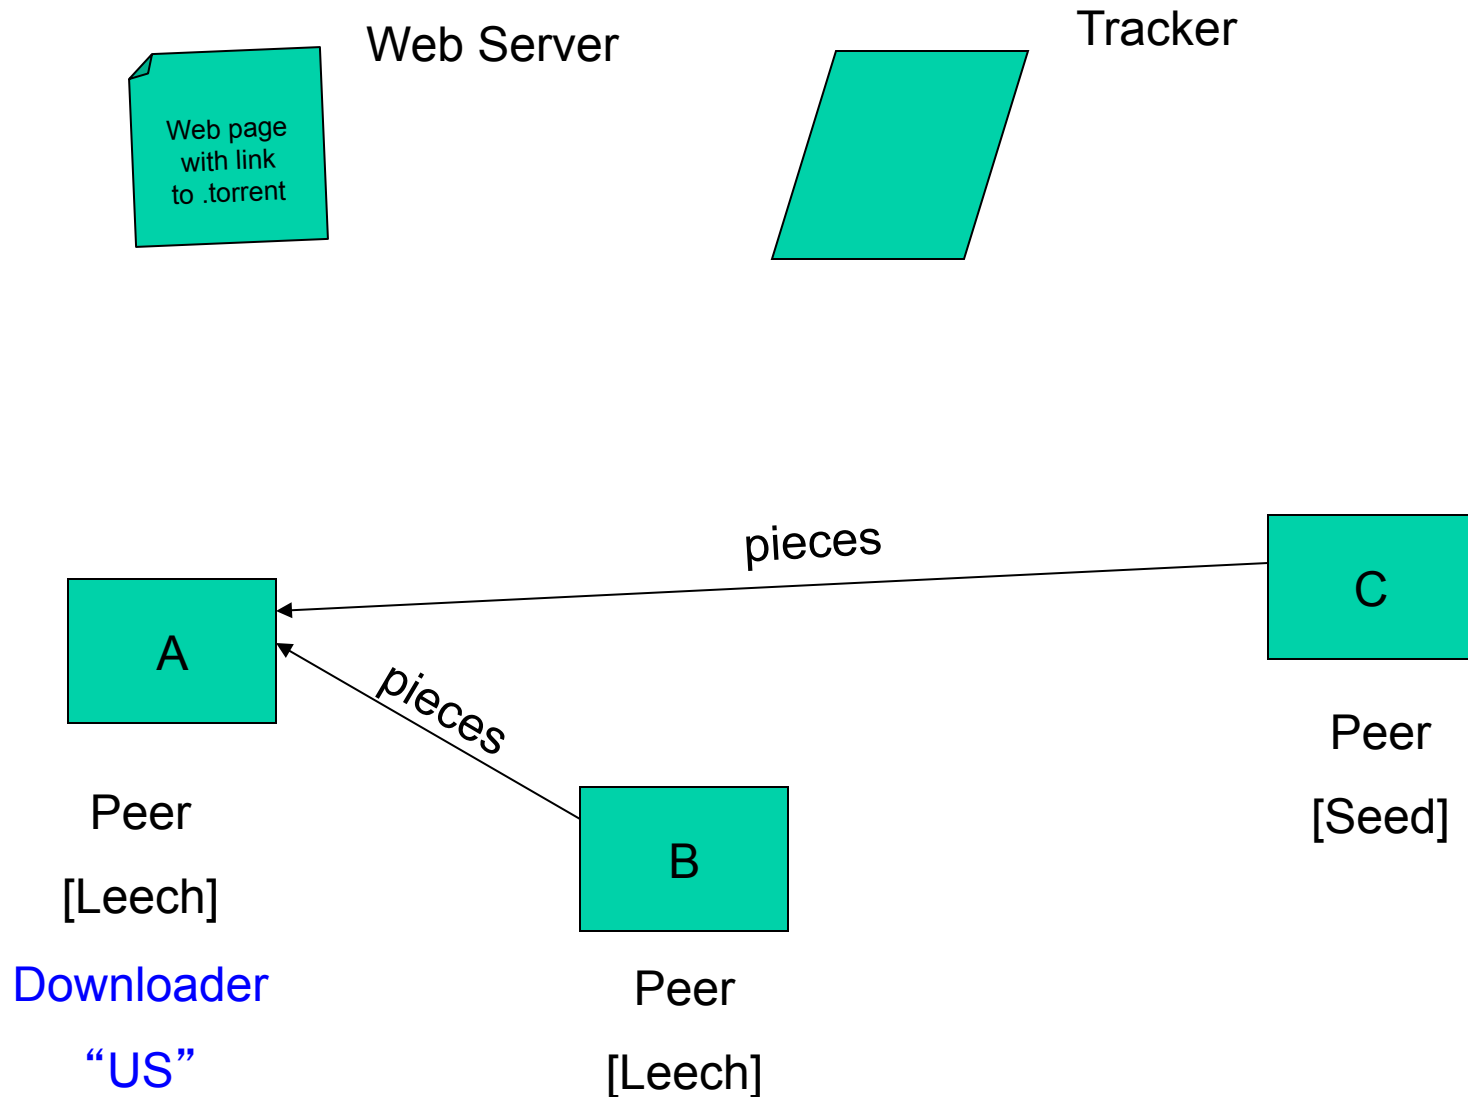
BitTorrent: Overall Architecture



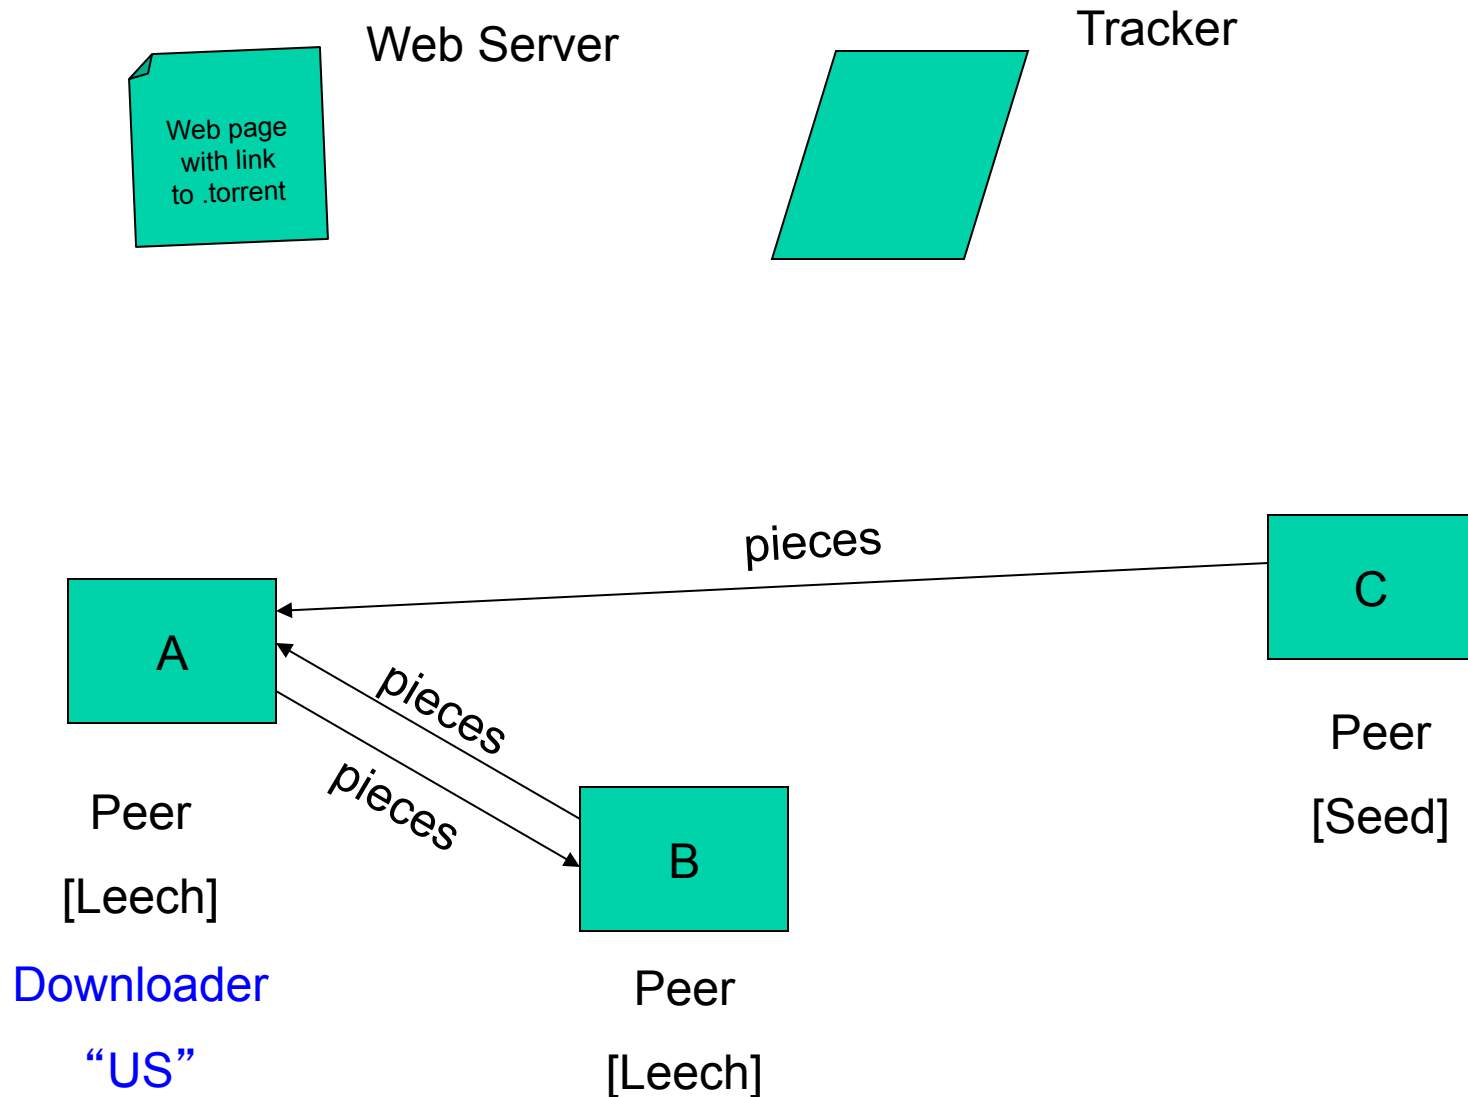
BitTorrent: Overall Architecture



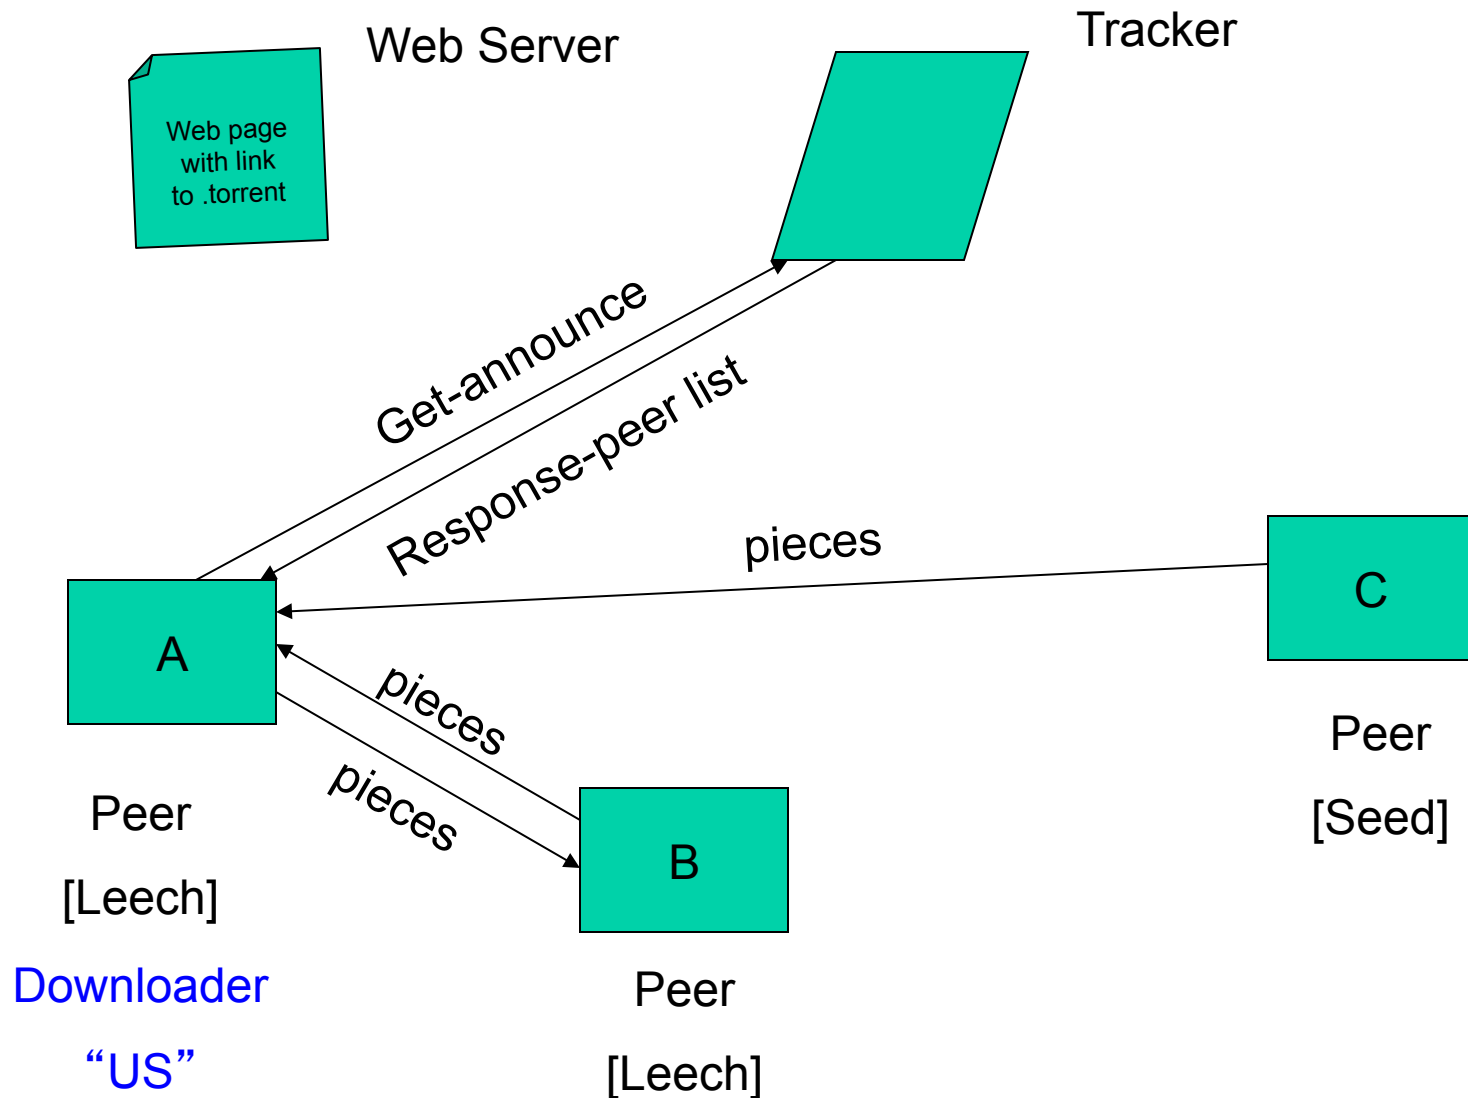
BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



Free-Riding Problem in P2P Networks

- Vast majority of users are free-riders
 - Most share no files and answer no queries
 - Others limit # of connections or upload speed
- A few “peers” essentially act as servers
 - A few individuals contributing to the public good
 - Making them hubs that basically act as a server
- How BitTorrent prevents free riding
 - Allow the fastest peers to download from you
 - Occasionally let some free loaders download

Conclusions

- Peer-to-peer networks
 - Nodes are end hosts
 - Primarily for file sharing, and recently telephony
 - Centralized directory (Napster), query flooding (Gnutella), super-nodes (KaZaA), and distributed downloading and anti-free-loading (BitTorrent)
- Great example of how change can happen so quickly in application-level protocols